

Neural Particle Method: Experimental Record

Synthetic-data setup, machinery, and the benchmark suite

Osian Shelley

8 September 2026

Abstract

This note records, in one place, the synthetic markets and dynamics used throughout the project, the simulation, estimation and scoring machinery, and the specification of the clean benchmark suite agreed on 8 September 2026 (Section 4), run on 9 September 2026; its tables are generated from the MLflow store. The suite follows the discovery, in the diagnostics of 8 September, that all one-pass schemes were fitting the conditional expectation of the raw Euler variance rather than of its positive part, which the dynamics actually use (Section 3.4); the diagnostic runs themselves live in the MLflow experiments `mechanism`, `bump_floor` and `bump_floor_nn`.

Contents

1	Conventions	1
2	Synthetic data	2
2.1	Target surfaces: SSVI power law	2
2.2	Variance dynamics	2
2.3	Quotes, bump, and scoring	2
3	Machinery	3
3.1	Simulation	3
3.2	One-pass schemes and their estimators	3
3.3	Implicit loop	4
3.4	The variance-target discrepancy	4
4	Benchmark suite	4
4.1	Common setup	4
4.2	Cold suite	5
4.3	What a lag is	5
4.4	Lagged suite	6
4.5	Per-scenario numbers	8
4.6	Reproducibility	10
5	Follow-up studies: the estimator, the body, and the online budget	10
5.1	Estimator quality against the exact conditional expectation	10
5.2	Training the body	10
5.3	The online budget	11
5.4	What the studies say	11

1 Conventions

Explicit denotes the one-pass, per-slice calibration with a per-slice neural network. Nadaraya–Watson (NW) is the classical particle method: the same one-pass structure with kernel regression. Other per-slice estimators are named by their estimator (RKHS, spline, ridge). *Implicit* denotes the damped outer iteration with a single network in (t, x) . All errors are implied-volatility root-mean-square errors in volatility basis points against the target surface, unless stated otherwise. “Pooled” means over all quoted strikes and maturities together; “wings” means the quoted strikes with $|k| > 0.25$; per-maturity columns restrict to one maturity.

2 Synthetic data

2.1 Target surfaces: SSVI power law

Each scenario’s market is an SSVI implied-variance surface in the power-law parameterisation. With log-moneyness $k = \log(K/S_0)$, $S_0 = 1$, and parameters $(\sigma_0, \eta, \gamma, \rho_s)$,

$$\theta_T = \sigma_0^2 T, \quad \varphi(\theta_T) = \eta \theta_T^{-\gamma}, \quad w(k, T) = \frac{1}{2} \theta_T \left(1 + \rho_s \varphi k + \sqrt{(\varphi k + \rho_s)^2 + 1 - \rho_s^2} \right), \quad (1)$$

where w is total implied variance and the implied volatility is $\sqrt{w(k, T)/T}$. Parameters are drawn independently per scenario from

$$\sigma_0 \sim U(0.12, 0.35), \quad \eta \sim U(0.5, 2.0), \quad \gamma \sim U(0.3, 0.5), \quad \rho_s \sim U(-0.85, -0.10),$$

resampled until the sufficient no-arbitrage conditions $0 < \gamma \leq \frac{1}{2}$, $|\rho_s| < 1$, $\eta(1 + |\rho_s|) \leq 2$ hold. Scenario s_i uses the seed $1000 + i$ (`bench/scenarios.py`, `make_registry`).

The Dupire local volatility is analytic from Equation (1): with w_k, w_{kk}, w_T the derivatives of w ,

$$\sigma_{\text{Dup}}^2(T, k) = \frac{w_T}{1 - \frac{k}{w} w_k + \frac{1}{4} \left(-\frac{1}{4} - \frac{1}{w} + \frac{k^2}{w^2} \right) w_k^2 + \frac{1}{2} w_{kk}}, \quad (2)$$

with the denominator floored at 0.05, the local variance clipped to $[10^{-8}, 9]$, and the time argument clipped to $[t_{\min}, T]$ with $t_{\min} = 0.004$ (`market/local_vol.py`). These clipping constants are part of the calibrated behaviour and are never changed.

2.2 Variance dynamics

The calibrated model is Heston-type LSV in log-spot $X = \log S$,

$$dX_t = -\frac{1}{2} L^2(t, X_t) V_t dt + L(t, X_t) \sqrt{V_t} dW_t, \quad dV_t = \kappa(\theta - V_t) dt + \xi \sqrt{V_t} dB_t, \quad d\langle W, B \rangle_t = \rho dt. \quad (3)$$

For scenario s_i the dynamics parameters are drawn from the same seed as the surface: $\kappa \in \{1, 2, 3\}$, $\xi \in \{0.3, 0.6, 1.0\}$, $\rho \in \{-0.7, -0.3\}$, each uniform over its set, and $\theta = v_0 = \sigma_0^2$ so that the variance level matches the surface’s short-end level. The Feller ratio $2\kappa\theta/\xi^2$ therefore ranges from 0.05 to 4.3 across the grid. Table 1 lists all twenty scenarios. Horizon $T = 2$ years; spot $S_0 = 1$; no rates or dividends.

Two further families exist in the registry and were not used today: six variants of s01 with (ξ, ρ) swept over the grid (`fig3_registry`), and three Heston-market scenarios whose target surface is generated by a Heston model with parameters different from the calibrated dynamics (`li_simple`, `li_complex`, `bayer`; horizon 1 year; Dupire surface from prices on a 50×121 grid).

2.3 Quotes, bump, and scoring

Quoted strikes are 13 log-strikes equally spaced in log between 0.6 and 1.6; maturities are 0.25, 0.5, 1 and 2 years. The wings are the strikes with $|k| > 0.25$ (six of thirteen). The warm-start

sid	SSVI surface				dynamics			Feller
	σ_0	η	γ	ρ_s	κ	ξ	ρ	
s01	0.261	0.52	0.34	-0.21	1	0.3	-0.7	1.51
s02	0.208	1.04	0.45	-0.56	3	0.6	-0.7	0.72
s03	0.163	0.89	0.40	-0.35	3	1.0	-0.7	0.16
s04	0.120	0.70	0.33	-0.56	2	0.6	-0.3	0.16
s05	0.285	1.01	0.49	-0.72	2	1.0	-0.3	0.32
s06	0.304	1.00	0.38	-0.58	3	1.0	-0.3	0.55
s07	0.202	0.56	0.46	-0.49	3	1.0	-0.3	0.25
s08	0.158	1.72	0.34	-0.11	1	1.0	-0.7	0.05
s09	0.133	0.99	0.31	-0.64	3	1.0	-0.3	0.11
s10	0.142	1.30	0.48	-0.11	3	1.0	-0.7	0.12
s11	0.144	1.10	0.38	-0.11	2	1.0	-0.7	0.08
s12	0.191	1.01	0.33	-0.22	3	1.0	-0.7	0.22
s13	0.303	0.62	0.43	-0.51	2	1.0	-0.3	0.37
s14	0.220	1.38	0.48	-0.11	1	0.6	-0.3	0.27
s15	0.127	0.72	0.37	-0.33	3	0.3	-0.3	1.08
s16	0.172	1.16	0.37	-0.42	1	0.6	-0.7	0.17
s17	0.136	1.13	0.42	-0.50	1	0.3	-0.3	0.41
s18	0.257	1.18	0.47	-0.24	3	0.6	-0.3	1.10
s19	0.170	1.09	0.39	-0.23	2	0.3	-0.7	1.28
s20	0.253	0.93	0.32	-0.32	3	0.3	-0.7	4.27

Table 1: The twenty synthetic scenarios: SSVI surface parameters, variance dynamics, and the Feller ratio $2\kappa\theta/\xi^2$.

experiments bump the surface to $\sigma'_0 = \sigma_0 + 0.01$, $\eta' = 0.95\eta$, $\rho'_s = \min(\rho_s + 0.03, -0.05)$, a level rise of one volatility point with curvature and skew nudged inside the no-arbitrage region.

Scoring reprices vanilla calls under the calibrated leverage field on fresh paths: 500 000 particles, 200 Euler steps over $[0, 2]$, seed 10 000 plus a strategy-specific offset, call price as the sample mean of the payoff, Black–Scholes inversion at the grid-snapped maturity, and the exact SSVI implied volatility at the same snapped time as the target (`pricing/reprice.py`, `pricing/metrics.py`). The bump experiments use 300 000 repricing paths. Errors are reported as RMSE in volatility bp, pooled, wings-only, and per maturity.

3 Machinery

3.1 Simulation

Calibration and repricing share one stepper. With $\Delta = T/K$, $v^+ = \max(v, 0)$, $z_1 = \rho z_b + \sqrt{1 - \rho^2} z_p$ and independent standard normals z_b, z_p ,

$$x_{k+1} = x_k - \frac{1}{2}L_k(x_k)^2 v_k^+ \Delta + L_k(x_k) \sqrt{v_k^+} \sqrt{\Delta} z_1, \quad v_{k+1} = v_k + \kappa(\theta - v_k^+) \Delta + \xi \sqrt{v_k^+} \sqrt{\Delta} z_b. \quad (4)$$

This is full truncation: the raw v may be negative and is carried, but only v^+ enters any coefficient. Calibration uses $K = 50$ steps unless stated; repricing always uses 200. The leverage is stored per slice on a grid in log-spot and evaluated by linear interpolation, constant beyond the grid ends; the slice in force at time t is the last grid time $\leq t$. The explicit and NW passes use a per-slice grid at the 0.1% to 99.9% quantiles of the current cloud (101 points); the implicit loop uses a fixed grid of 81 points from log 0.4 to log 2.2. The leverage is capped at $L_{\max} = 4$ and the fitted conditional expectation is floored at 10^{-4} .

3.2 One-pass schemes and their estimators

At slice k , fit $f_k(x) \approx \mathbb{E}[V_{t_k} \mid X_{t_k} = x]$ on a subsample of 30 000 particles from the current cloud (the paper’s setting; the benchmark suite of Section 4 fits on the full cloud), set $L_k = \text{clip}(\sigma_{\text{Dup}}(t_k, \cdot) / \sqrt{\max(f_k, 10^{-4})}, 0, L_{\text{max}})$, and advance the cloud. Slice 0 uses $f_0 \equiv v_0$. Estimators:

- **NW**: Gaussian Nadaraya–Watson with Silverman bandwidth $h = 1.06 \hat{\sigma}_x n^{-1/5}$ on the subsample.
- **Explicit (per-slice network)**: a $1 \rightarrow 64 \rightarrow 64 \rightarrow 1$ SiLU network with softplus output scaled by 0.04, input scaled by 0.3, trained by Adam at learning rate 10^{-2} on the unweighted squared loss; 400 steps on the first fitted slice and 120 steps on each later slice, warm-started from the previous slice’s weights.
- **RKHS**: kernel ridge with a Gaussian kernel of variance 0.1 in log-spot, 100 centres at the $j \cdot 100/101$ percentiles of the cloud, regularisation $\lambda = 10^{-9}$ (Bayer et al.’s published choices).
- **Ridge head on a frozen body**: the implicit network’s body (below) as a 65-dimensional feature map, with a per-slice ridge readout, $\lambda = 10^{-3}$ per sample, shrunk toward the previous slice’s coefficients.

3.3 Implicit loop

A single network $f_\theta(t, x)$, $2 \rightarrow 64 \rightarrow 64 \rightarrow 1$ SiLU with softplus output scaled by 0.04, inputs $(t/T, x/0.3)$. Starting from the explicit output as L^0 , each outer iteration simulates a full cloud under L^n with fresh random numbers, draws a pooled subsample of 60 000 points across all slices, trains θ for 300 Adam steps at learning rate 10^{-2} with the network and optimiser state carried over from the previous iteration, reads $\Phi(L^n)_k = \text{clip}(\sigma_{\text{Dup}}(t_k) / \sqrt{\max(f_\theta(t_k, \cdot), 10^{-4})}, 0, L_{\text{max}})$ on the fixed grid, and sets $L^{n+1} = (1 - \alpha)L^n + \alpha\Phi(L^n)$ with $\alpha = \frac{1}{2}$. The paper uses six iterations. The pooled data are truncated, v^+ , before training.

3.4 The variance-target discrepancy

The one-pass schemes hand the estimator the raw Euler v ; the implicit loop hands it v^+ . Since only v^+ drives the dynamics Equation (4), the quantity the leverage must divide by is $\mathbb{E}[V^+ \mid X]$. Under Feller violation a sizeable fraction of particles has $v < 0$ after a coarse step (on s11, $v_0 = 0.02$, $\xi = 1$, $\Delta = 0.04$: the one-step standard deviation of v is 0.028), so $\mathbb{E}[v \mid X] < \mathbb{E}[v^+ \mid X]$ by the mean negative excursion, the denominator is too small, and the leverage too large. The effect is independent of the particle count and shrinks with the step size. “Floor on” below means the estimator is given v^+ (`ExplicitConfig.fit_v_floor`); the default remains off so that the golden replays of the paper’s runs are unchanged.

4 Benchmark suite

This section specifies the clean benchmark agreed on 8 September 2026 and run on 9 September 2026 with the variance floor of Section 3.4 as the package default. Every cell of the tables below is one MLflow run, keyed by method, scenario, seed and (for the lagged suite) lag type and offline particle count, in experiments `suite_cold`, `suite_offline`, `suite_lagged` and `pde_reference`. Every trained object is stored as a run artifact (network weights as a `torch` state dictionary, per-slice estimators and leverage fields as pickles and JSON) and can be reloaded by run id. The tables are generated from the store by `nparticle suite tables`.

4.1 Common setup

- **Scenarios**. The 20 SSVI scenarios of Table 1 and the three Heston-market scenarios used in the literature comparisons (`li_simple`, `li_complex`, `bayer`; a Heston market surface, a different

Heston variance process as the calibrated model, one-year maturity, log-strikes on $[-0.9, 0.9]$. The Heston block is reported separately because it scores a single one-year maturity (the same 13 quoted strikes).

- **Time grid.** 200 Euler steps for calibration and 200 for repricing, $\Delta = 0.01$ on $[0, 2]$, so that every method sits at its step-count plateau (the step-count diagnostics of 8 September 2026, MLflow experiment `mechanism`) and the comparison is about the estimator alone.
- **Particles.** 100 000 particles for every cold row and for every online stage; every per-slice fit and every online head is fitted on the full cloud, not on the paper’s 30 000 subsample. Offline bodies at 200 000 and 500 000 particles, fitted on 100 000-particle subsamples.
- **Seeds.** Two seeds (0, 1) for every cold row and every online stage; one seed for each offline body. The reprice seed is 10 000 plus the calibration seed.
- **Floor.** All one-pass schemes fit $\mathbb{E}[V^+ | X]$; `fit_v_floor` becomes the default in this work and the tiny golden replays are regenerated in the same commit.
- **Implicit configuration.** Warm start from the explicit network field, 30 damped iterations at $\alpha = \frac{1}{2}$ with fresh noise per iteration, persistent network and optimiser (the best practical setting found in the convergence diagnostics of 8 September 2026, experiment `mechanism`).
- **Scoring.** As in Section 2.3: 500 000 repricing paths, Black–Scholes inversion at the snapped maturity, exact SSVI (or semi-analytic Heston) implied volatility as the target. Metrics per run: pooled and wings MAE and RMSE in volatility bp, both per maturity, the count of failed inversions, and the leverage RMSE against the PDE reference of the scenario. Errors are reported in implied-vol basis points (the headline) and, in the companion tables, as absolute option-price errors in basis points of spot, recovered exactly from the implied-vol errors through the Black–Scholes formula.
- **Latency.** Wall-clock seconds of the calibration only (repricing excluded). For the lagged suite, only the online stage is timed.
- **PDE reference.** The Fokker–Planck reference leverage (`reference/pde.py`) is solved per scenario at 200 steps and enters twice: repriced under the same Monte Carlo as every other row, which is the attainable floor of the scoring under this repricer (Gyöngy), and as the reference for the leverage-error column. Its repricing is repeated under a second reprice seed to measure the Monte Carlo noise floor of the score itself.

4.2 Cold suite

One calibration from scratch per method, scenario and seed at 100 000 particles. Methods, with their estimator of $\mathbb{E}[V^+ | X]$ per slice:

1. **NW:** Gaussian Nadaraya–Watson, Silverman bandwidth (Section 3.2).
2. **Explicit NN:** per-slice network, warm-started across slices (Section 3.2).
3. **Implicit NN:** damped global-network loop (Section 3.3, configuration above).
4. **RKHS:** Gaussian kernel ridge with 100 quantile centres, Bayer et al.’s published knobs (Section 3.2).
5. **Spline:** penalised cubic B-spline, 25 quantile knots, $\lambda = 1$.
6. **Literature block.** GHL kernel: Nadaraya–Watson in spot with the quartic kernel and the Guyon–Henry–Labordère rule-of-thumb bandwidth $h(t) = 1.5 S_0 \sigma_{LV}(S_0, t) \sqrt{\max(t, 0.25)} N^{-1/5}$. Bins: piecewise-constant regression on 20 equal-frequency buckets (van der Stoep et al.). Muguruza: the kernel-free conditional Monte Carlo of Corollary 4.1, each particle contributing its one-step Gaussian conditional law of log-spot given the variance path. PU-RBF: Hakala’s partition-of-unity radial basis functions, 40 centres, 5-nearest-neighbour widths, regulariser 0.2.

7. **PDE**: the reference leverage repriced by Monte Carlo (attainable floor).

Table 2: Cold suite, SSVI scenarios: mean and median over the 20 scenarios and two seeds. The PDE row is the attainable floor of the scoring; its seed-to-seed spread (median over scenarios of the absolute difference in pooled MAE between the two reprice seeds) is 6.6 bp, the Monte Carlo noise floor of the score.

method	pooled MAE (bp)		pooled RMSE (bp)		wings MAE (bp)		lev. RMSE	latency (s)
	mean	median	mean	median	mean	median	median	median
NW	46	33	87	67	70	48	0.093	8.0
Explicit NN	57	36	112	64	86	49	0.089	1042.7
Implicit NN	58	37	110	73	88	47	0.395	1298.6
RKHS	81	70	131	110	113	99	0.224	12.4
Spline	47	34	90	64	73	46	0.098	2.7
GHL kernel	45	32	88	62	70	45	0.106	8.5
Bins	66	61	120	99	105	84	0.210	1.5
Muguruza	154	139	188	171	147	136	0.161	10.6
PU-RBF	57	40	112	66	90	53	0.130	2.5
PDE (attainable floor)	42	27	85	59	62	35	0.000	0.0

Table 3: Cold suite, SSVI scenarios, the same runs in price space: absolute option-price error in basis points of spot (spot is one), mean over scenarios and seeds. Relative price error is not shown; in the far wings the target prices are of order 10^{-6} of spot.

method	pooled mean	pooled median	near-money	wings	$T = 2$	worst quote
NW	2.7	2.4	3.6	1.8	4.3	9.3
Explicit NN	2.6	2.1	3.4	2.0	3.7	8.6
Implicit NN	3.3	2.7	4.6	2.1	5.5	12.2
RKHS	8.5	8.1	11.6	5.7	19.0	32.3
Spline	2.5	2.1	3.3	1.7	3.8	8.2
GHL kernel	2.5	2.2	3.4	1.7	4.0	8.6
Bins	2.7	2.3	3.3	2.2	4.2	8.5
Muguruza	29.1	35.0	47.1	12.3	62.7	122.9
PU-RBF	2.8	2.5	3.4	2.2	4.4	9.2
PDE (attainable floor)	2.2	1.7	3.2	1.3	3.1	8.1

4.3 What a lag is

An overnight calibration is done on surface S_0 ; the intraday market is S_1 , and the question is how much of the overnight work survives. Because the synthetic surfaces are parametric, a lag is a move in parameter space and never a bump of quoted points, so arbitrage cannot be introduced by accident. Two lags are run for every cell.

Surface lag. For SSVI, the existing bump: $\sigma'_0 = \sigma_0 + 0.01$, $\eta' = 0.95\eta$, $\rho'_s = \min(\rho_s + 0.03, -0.05)$, shrunk by quarters until the Gatheral–Jacquier sufficient conditions ($\eta(1 + |\rho_s|) \leq 2$, $0 < \gamma \leq \frac{1}{2}$) hold; calendar no-arbitrage is automatic since $\theta_T = \sigma_0^2 T$ is increasing. For a Heston market, the matching move in vol terms: $\sqrt{v_0}$ and $\sqrt{\theta}$ up one volatility point, $\rho \rightarrow \min(\rho + 0.03, -0.05)$, $\xi \rightarrow 0.95\xi$; the surface is model-generated and arbitrage-free for any valid parameters.

Surface plus spot lag. The surface lag together with a sticky-strike spot move of +2%, $S_1 = S_0 e^\delta$, $\delta = \log 1.02$. Under sticky-strike the implied volatility at each absolute strike is unchanged, so in log-moneyness relative to the new spot the surface is translated: $\sigma'_{IV}(k, T) = \sigma_{IV}(k + \delta, T)$. A translation in k preserves Durrleman’s butterfly condition (it involves only w and its k -derivatives at a point) and calendar monotonicity, so the shifted surface is arbitrage-free. Call prices at fixed strike change with the spot, so the Dupire local volatility is that of the translated surface, $\tilde{w}(k) = w(k + \delta)$ with k measured from the new forward; the particles also start at the new spot. The stored denominators $f(t, x) = \mathbb{E}[V_t | X_t = x]$ were fitted for a cloud started at S_0 , and a conditional expectation given

Table 4: Cold suite, Heston-market scenarios (one-year maturity, same 13 quotes): mean and median over the three scenarios and two seeds..

method	pooled MAE (bp)		pooled RMSE (bp)		wings MAE (bp)		lev. RMSE	latency (s)
	mean	median	mean	median	mean	median	median	median
NW	193	61	218	66	125	64	0.145	7.9
Explicit NN	189	55	214	60	120	57	0.143	1044.4
Implicit NN	179	51	206	56	107	51	0.528	1281.5
RKHS	209	67	242	74	147	77	0.241	12.2
Spline	193	57	218	62	124	60	0.165	2.7
GHL kernel	194	56	218	60	126	58	0.145	7.8
Bins	186	61	218	73	117	71	0.345	1.5
Muguruza	381	460	408	494	293	398	0.753	8.2
PU-RBF	193	55	218	60	124	57	0.145	2.3
PDE (attainable floor)	189	40	211	44	123	43	0.000	0.0

Table 5: Cold suite, Heston-market scenarios, the same runs in price space: absolute option-price error in basis points of spot (spot is one), mean over scenarios and seeds. Relative price error is not shown; in the far wings the target prices are of order 10^{-6} of spot.

method	pooled mean	pooled median	near-money	wings	worst quote
NW	50.5	13.2	89.3	17.2	135.1
Explicit NN	50.0	12.2	88.4	17.1	134.2
Implicit NN	48.4	11.7	86.4	15.9	131.8
RKHS	52.4	13.0	91.5	18.9	140.5
Spline	50.7	12.4	89.5	17.4	135.6
GHL kernel	50.6	12.4	89.4	17.4	135.5
Bins	48.6	11.4	86.9	15.8	132.3
Muguruza	93.0	110.5	157.1	38.0	223.2
PU-RBF	50.8	12.4	89.6	17.5	135.7
PDE (attainable floor)	48.9	8.5	86.7	16.5	130.6

spot depends on the starting point: with $\rho < 0$ a particle sitting 2% above the old start had its variance drift down, so the stale f at the new spot is below v_0 , the stale leverage is too high near the money, most at short maturities. Sticky-delta, by contrast, leaves the surface in moneyness unchanged and the whole calibration is scale invariant, so a sticky-delta spot move is invisible to every method; the surface-only lag already covers it. A 2% daily move is one to one-and-a-half standard deviations at 20–30% volatility.

4.4 Lagged suite

Offline stage (experiment `suite_offline`, one run per scenario, body type and particle count): a full calibration on S_0 at 200 000 or 500 000 particles, 200 steps, seed 0: either the explicit NN pass (the per-slice networks and the per-slice denominators f_k are saved) or the implicit loop (the global network is saved). These runs are cached and reused by every online cell below.

Online stage (experiment `suite_lagged`, one run per scenario, method, offline size, lag and online seed): given the intraday surface S_1 and 100 000 fresh particles, build the intraday leverage and reprice. Only this stage is timed. Methods:

1. **Explicit NN, stale f , fresh Dupire:** $L_k = \text{clip}(\sigma_{\text{Dup}}^{S_1}(t_k, \cdot) / \sqrt{\max(f_k, 10^{-4})}, 0, L_{\max})$ with the offline per-slice f_k . Costs nothing online beyond the repricing.
2. **Implicit NN, stale f , fresh Dupire:** the same with $f = f_\theta(t, \cdot)$ from the offline global network.
3. **Explicit NN with RKHS head:** simulate the 100 000 online particles under the stale- f leverage of item 1 with fresh Dupire, slice by slice; at each slice fit a Gaussian kernel ridge

(kernel width 0.5 times the cloud’s standard deviation, ridge 10^{-3}) to the residual $v^+ - f_k(x)$ on the online cloud, and use $\max(f_k + \text{residual}, \frac{1}{4}f_k)$ as the denominator for the next step; the floor keeps the head from driving the leverage to its cap on Feller-violating scenarios.

4. **Implicit NN with RKHS head:** the same residual head on the global network’s f_θ .
5. **Explicit NN with ridge head:** the per-slice network body frozen as a feature map (64 features plus a constant) with its last layer refitted by ridge on the online cloud, shrunk toward the offline readout. The body’s output is $\text{softplus}(Aw) \cdot 0.04$, and the refit minimises $\sum_i (\text{softplus}(A_i w) \cdot 0.04 - v_i^+)^2 + \lambda_{\text{eff}} \|w - w_0\|^2$ on the raw variance scale by Gauss–Newton to convergence (at most 20 steps) from the offline last layer w_0 , with $\lambda_{\text{eff}} = \lambda$ times the mean diagonal of the Gauss–Newton curvature at w_0 and $\lambda = 10^{-3}$, so the penalty is a fixed fraction of the data curvature whatever the output scale and cloud size. The denominator is $\text{softplus}(Aw) \cdot 0.04$ at the converged w . When the online data agree with the offline fit the first step is zero and $w = w_0$.
6. **Implicit NN with ridge head:** the same on the global body in (t, x) .
7. **Explicit NN with spline head and Implicit NN with spline head:** the same residual correction with the penalised cubic B-spline of Section 3.2 (25 quantile knots, $\lambda = 1$) in place of the kernel ridge, with the same floor.

The online reference rows are NW re-solved from scratch on S_1 at 100 000 particles and 200 steps (the best intraday strategy in the floor-corrected bump runs, experiment `bump_floor`), and the stale leverage with no action.

4.5 Per-scenario numbers

Cold suite, SSVI scenarios, pooled MAE (bp)

sid	NW	Explicit NN	Implicit NN	RKHS	Spline	GHL kernel	Bins	Muguruza	PU-RBF
s01	15	13	16	17	14	14	23	13	14
s02	12	15	20	38	14	14	20	30	27
s03	61	107	58	160	65	63	101	238	115
s04	106	104	149	120	111	105	144	187	107
s05	50	43	39	59	45	44	62	137	44
s06	26	24	24	104	27	25	34	82	26
s07	47	45	49	69	51	46	79	185	47
s08	75	124	81	130	78	77	112	435	110
s09	136	127	172	214	136	136	160	260	138
s10	29	36	80	94	30	29	61	344	36
s11	71	204	145	160	80	72	127	333	134
s12	53	52	73	83	57	51	85	229	86
s13	27	26	22	60	28	27	36	137	27
s14	24	26	32	32	24	25	21	140	26
s15	71	71	69	71	71	72	77	72	72
s16	36	37	35	109	39	34	69	163	62
s17	26	26	28	37	27	26	43	38	25
s18	18	19	30	32	18	19	15	28	19
s19	22	22	24	21	23	22	35	23	24
s20	11	9	11	11	11	10	19	10	9

Cold suite, Heston scenarios, pooled MAE (bp)

Table 6: Lagged suite, SSVI scenarios: mean over the 20 scenarios and two online seeds; latency is the online stage only. One block per lag type..

method	offline N	pooled MAE	pooled RMSE	wings MAE	$T = 0.25$ MAE	online latency (s)
<i>Surface lag</i>						
Stale leverage, no action	–	99	133	105	141	0.0
NW re-solve on S_1	–	44	82	64	96	8.5
Explicit NN, stale f , fresh Dupire	200k	55	101	78	121	0.0
Explicit NN, stale f , fresh Dupire	500k	55	102	77	123	0.0
Implicit NN, stale f , fresh Dupire	200k	62	113	91	128	0.0
Implicit NN, stale f , fresh Dupire	500k	67	114	94	137	0.0
Explicit NN + RKHS head	200k	46	90	69	108	17.0
Explicit NN + RKHS head	500k	46	91	69	109	17.1
Implicit NN + RKHS head	200k	47	93	72	108	17.1
Implicit NN + RKHS head	500k	44	85	65	102	17.0
Explicit NN + ridge head	200k	47	93	71	113	30.1
Explicit NN + ridge head	500k	48	95	72	117	28.4
Implicit NN + ridge head	200k	50	98	75	118	33.1
Implicit NN + ridge head	500k	44	85	65	103	32.6
Explicit NN + spline head	200k	43	83	63	99	7.5
Explicit NN + spline head	500k	43	82	63	98	7.3
Implicit NN + spline head	200k	46	89	70	103	7.2
Implicit NN + spline head	500k	42	81	62	96	7.3
<i>Surface plus 2% sticky-strike spot lag</i>						
Stale leverage, no action	–	97	137	108	165	0.0
NW re-solve on S_1	–	44	83	65	98	8.5
Explicit NN, stale f , fresh Dupire	200k	97	139	114	198	0.0
Explicit NN, stale f , fresh Dupire	500k	97	141	115	200	0.0
Implicit NN, stale f , fresh Dupire	200k	98	143	121	192	0.0
Implicit NN, stale f , fresh Dupire	500k	97	140	119	184	0.0
Explicit NN + RKHS head	200k	46	90	70	110	16.9
Explicit NN + RKHS head	500k	46	92	70	111	17.1
Implicit NN + RKHS head	200k	49	94	74	110	17.1
Implicit NN + RKHS head	500k	45	87	68	101	17.0
Explicit NN + ridge head	200k	46	90	68	109	31.5
Explicit NN + ridge head	500k	48	93	71	114	29.3
Implicit NN + ridge head	200k	49	94	72	110	34.1
Implicit NN + ridge head	500k	46	90	68	109	32.8
Explicit NN + spline head	200k	44	84	65	100	7.5
Explicit NN + spline head	500k	44	84	65	100	7.3
Implicit NN + spline head	200k	49	92	75	104	7.3
Implicit NN + spline head	500k	56	100	87	108	7.3

sid	NW	Explicit NN	Implicit NN	RKHS	Spline	GHL kernel	Bins	Muguruza	PU-RBF
bayer	507	502	478	521	511	515	481	661	513
li_complex	61	55	51	67	57	56	61	460	55
li_simple	12	10	8	39	12	11	15	22	11

Lagged suite, surface lag, SSVI scenarios, pooled MAE (bp)

[illegible]

Lagged suite, surface plus 2% spot lag, SSVI scenarios, pooled MAE (bp)

	2000		2001		2002		2003		2004		2005		2006		2007		2008		2009		2010		2011		2012		2013		2014		2015		2016		2017		2018		2019		2020		2021		2022		2023		2024		2025		2026		2027		2028		2029		2030		2031		2032		2033		2034		2035		2036		2037		2038		2039		2040		2041		2042		2043		2044		2045		2046		2047		2048		2049		2050		2051		2052		2053		2054		2055		2056		2057		2058		2059		2060		2061		2062		2063		2064		2065		2066		2067		2068		2069		2070		2071		2072		2073		2074		2075		2076		2077		2078		2079		2080		2081		2082		2083		2084		2085		2086		2087		2088		2089		2090		2091		2092		2093		2094		2095		2096		2097		2098		2099		2100		2101		2102		2103		2104		2105		2106		2107		2108		2109		2110		2111		2112		2113		2114		2115		2116		2117		2118		2119		2120		2121		2122		2123		2124		2125		2126		2127		2128		2129		2130		2131		2132		2133		2134		2135		2136		2137		2138		2139		2140		2141		2142		2143		2144		2145		2146		2147		2148		2149		2150		2151		2152		2153		2154		2155		2156		2157		2158		2159		2160		2161		2162		2163		2164		2165		2166		2167		2168		2169		2170		2171		2172		2173		2174		2175		2176		2177		2178		2179		2180		2181		2182		2183		2184		2185		2186		2187		2188		2189		2190		2191		2192		2193		2194		2195		2196		2197		2198		2199		2200		2201		2202		2203		2204		2205		2206		2207		2208		2209		2210		2211		2212		2213		2214		2215		2216		2217		2218		2219		2220		2221		2222		2223		2224		2225		2226		2227		2228		2229		2230		2231		2232		2233		2234		2235		2236		2237		2238		2239		2240		2241		2242		2243		2244		2245		2246		2247		2248		2249		2250		2251		2252		2253		2254		2255		2256		2257		2258		2259		2260		2261		2262		2263		2264		2265		2266		2267		2268		2269		2270		2271		2272		2273		2274		2275		2276		2277		2278		2279		2280		2281		2282		2283		2284		2285		2286		2287		2288		2289		2290		2291		2292		2293		2294		2295		2296		2297		2298		2299		2300		2301		2302		2303		2304		2305		2306		2307		2308		2309		2310		2311		2312		2313		2314		2315		2316		2317		2318		2319		2320		2321		2322		2323		2324		2325		2326		2327		2328		2329		2330		2331		2332		2333		2334		2335		2336		2337		2338		2339		2340		2341		2342		2343		2344		2345		2346		2347		2348		2349		2350		2351		2352		2353		2354		2355		2356		2357		2358		2359		2360		2361		2362		2363		2364		2365		2366		2367		2368		2369		2370		2371		2372		2373		2374		2375		2376		2377		2378		2379		2380		2381		2382		2383		2384		2385		2386		2387		2388		2389		2390		2391		2392		2393		2394		2395		2396		2397		2398		2399		2400		2401		2402		2403		2404		2405		2406		2407		2408		2409		2410		2411		2412		2413		2414		2415		2416		2417		2418		2419		2420		2421		2422		2423		2424		2425		2426		2427		2428		2429		2430		2431		2432		2433		2434		2435		2436		2437		2438		2439		2440		2441		2442		2443		2444		2445		2446		2447		2448		2449		2450		2451		2452		2453		2454		2455		2456		2457		2458		2459		2460		2461		2462		2463		2464		2465		2466		2467		2468		2469		2470		2471		2472		2473		2474		2475		2476		2477		2478		2479		2480		2481		2482		2483		2484		2485		2486		2487		2488		2489		2490		2491		2492		2493		2494		2495		2496		2497		2498		2499		2500		2501		2502		2503		2504		2505		2506		2507		2508		2509		2510		2511		2512		2513		2514		2515		2516		2517		2518		2519		2520		2521		2522		2523		2524		2525		2526		2527		2528		2529		2530		2531		2532		2533		2534		2535		2536		2537		2538		2539		2540		2541		2542		2543		2544		2545		2546		2547		2548		2549		2550		2551		2552		2553		2554		2555		2556		2557		2558		2559		2560		2561		2562		2563		2564		2565		2566		2567		2568		2569		2570		2571		2572		2573		2574		2575		2576		2577		2578		2579		2580		2581		2582		2583		2584		2585		2586		2587		2588		2589		2590		2591		2592		2593		2594		2595		2596		2597		2598		2599		2600		2601		2602		2603		2604		2605		2606		2607		2608		2609		2610		2611		2612		2613		2614		2615		2616		2617		2618		2619		2620		2621		2622		2623		2624		2625		2626		2627		2628		2629		2630		2631		2632		2633		2634		2635		2636		2637		2638		2639		2640		2641		2642		2643		2644		2645		2646		2647		2648		2649		2650		2651		2652		2653		2654		2655		2656		2657		2658		2659		2660		2661		2662		2663		2664		2665		2666		2667		2668		2669		2670		2671		2672		2673		2674		2675		2676		2677		2678		2679		2680		2681		2682		2683		2684		2685		2686		2687		2688		2689		2690		2691		2692		2693		2694		2695		2696		2697		2698		2699		2700		2701		2702		2703		2704		2705		2706		2707		2708		2709		2710		2711		2712		2713		2714		2715		2716		2717		2718		2719		2720		2721		2722		2723		2724		2725		2726		2727		2728		2729		2730		2731		2732		2733		2734		2735		2736		2737		2738		2739		2740		2741		2742		2743		2744		2745		2746		2747		2748		2749		2750		2751		2752		2753		2754		2755		2756		2757		2758		2759		2760		2761		2762		2763		2764		2765		2766		2767		2768		2769		2770		2771		2772		2773		2774		2775		2776		2777		2778		2779		2780		2781		2782		2783		2784		2785		2786		2787		2788		2789		2790		2791		2792		2793		2794		2795		2796		2797		2798		2799		2800		2801		2802		2803		2804		2805		2806		2807		2808		2809		2810		2811		2812		2813		2814		2815		2816		2817		2818		2819		2820		2821		2822		2823		2824		2825		2826		2827		2828		2829		2830		2831		2832		2833		2834		2835		2836		2837		2838		2839		2840		2841		2842		2843		2844		2845		2846		2847		2848		2849		2850		2851		2852		2853		2854		2855		2856		2857		2858		2859		2860		2861		2862		2863		2864		2865		2866		2867		2868		2869		2870		2871		2872		2873		2874		2875		2876		2877		2878		2879		2880		2881		2882		2883		2884		2885		2886		2887		2888		2889		2890		2891		2892		2893		2894		2895		2896		2897		2898		2899		2900		2901		2902		2903		2904		2905		2906		2907		2908		2909		2910		2911		2912		2913		2914		2915		2916		2917		2918		2919		2920		2921		2922		2923		2924		2925		2926		2927		2928		2929		2930		2931		2932		2933		2934		2935		2936		2937		2938		2939		2940		2941		2942		2943		2944		2945		2946		2947		2948		2949		2950		2951		2952		2953		2954		2955		2956		2957		2958		2959		2960		2961		2962		2963		2964		2965		2966		2967		2968		2969		2970		2971		2972		2973		2974		2975		2976		2977		2978		2979		2980		2981		2982		2983		2984		2985		2986		2987		2988		2989		2990		2991		2992		2993		2994		2995		2996		2997		2998		2999		3000		3001		3002		3003		3004		3005		3006		3007		3008		3009		3010		3011		3012		3013		3014		3015		3016		3017		3018		3019		3020		3021		3022		3023		3024		3025		3026		3027		3028		3029		3030	
--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--	------	--

Lagged suite, surface lag, Heston scenarios, pooled MAE (bp)

[illegible]

Lagged suite, surface plus 2% spot lag, Heston scenarios, pooled MAE (bp)

[illegible]

4.6 Reproducibility

Each run records the git hash, the full configuration as parameters, the leverage field, the implied-volatility error grid, and the trained objects. A loader (`nparticle suite load <run_id>`) rebuilds the leverage field and, where present, the network or per-slice estimators from the artifacts alone, so any row of the tables can be re-scored or reused as an offline body without recomputation.

5 Follow-up studies: the estimator, the body, and the online budget

Three studies run on 10–11 September 2026 after the suite, on the six scenarios s01, s02, s05, s09, s11 and s16 (Feller ratios 1.5, 0.7, 0.3, 0.1, 0.08, 0.17). Every fit in these studies, and in the suite tables of Section 4, uses the full online cloud; the suite’s first pass, which fitted every per-slice estimator on the paper’s 30 000-particle subsample, is kept in the store under the experiment names `suite*_sub30k` and is quoted below only where the comparison is the point. Scripts: `scripts/estimator_grid.py`, `scripts/conv_study.py`, `scripts/budget_sweep.py`; tables from `scripts/study_tables.py`.

5.1 Estimator quality against the exact conditional expectation

Setup. For s02 (an easy surface) and s11 (Feller 0.08), a cloud of 500 000 particles is simulated under the PDE reference leverage of Section 4.1, and at the slices $t = 0.25$ and $t = 1$ each estimator is fitted on (x, v^+) and compared with the reference’s own conditional expectation f_{ref} on the resolved grid. The error is $|\hat{f} - f_{\text{ref}}|/f_{\text{ref}}$, averaged over the wings ($|x| > 0.25$) or the centre. Networks are the per-slice architecture of Section 3.2 with the hidden width and depth varied, trained from a cold start on 100 000 particles; “full batch” is Adam on the whole sample at learning rate 10^{-2} , “minibatch” is Adam on batches of 8192 at 10^{-3} . NW is fitted on 30 000, 100 000 and 500 000 particles of the same cloud.

Conclusions. (i) The 30 000 subsample was costing real accuracy in the wings: NW’s wing error halves from 30k to 500k particles. (ii) On the easy surface the lever is training, not architecture: width 256 and depth 3 gain nothing, while the same 64×2 network trained with 20 000 minibatch steps reaches 0.6–1.7% in the wings, better than NW on five times the particles, in under two minutes per slice; the paper’s 120 full-batch steps leave it at 2.6–6% from a cold start (the warm start across slices in the real pass softens this). (iii) On the Feller-violating scenario every estimator, NW on 500k included, is 10–15% off the reference even at the money, and more particles do not help: the Euler cloud with full truncation at $\Delta = 0.01$ is not the continuous-time density, so the reference cannot judge estimators there. (iv) Every 2-layer network collapses to 54% wing error at $t = 0.25$ on that scenario, at every budget, while the 3-layer networks do not: depth buys robustness where the variance has an atom at zero, not accuracy on easy surfaces.

5.2 Training the body

Setup. The offline bodies of Section 4.4 at 500 000 particles were retrained on the six scenarios under two training budgets and refitted on the full online cloud: “short” is the paper’s schedule (400 Adam steps on the first slice, 120 on each later slice; 300 steps per outer iteration for the implicit loop), “long” is 2000 and 1000 steps per slice (1000 per iteration on a 200 000-point pool for the implicit loop). Table 11 scores each body as a cold calibrator on its own overnight surface; Table 12 runs the

lagged suite on it with every fit on the full 100 000-particle online cloud, alongside the first pass’s 30k-subsample rows.

Conclusions. (i) Training the explicit body ten times longer turns it from the weakest cold calibrator into the best one on these scenarios (64 $\rightarrow$ 50 pooled against NW’s 53), at 2.5 hours per body with full-batch training; the minibatch schedule of Section 5.1 should reach the same quality in a fraction of that and is untested as a body. (ii) The implicit loop gets worse when trained harder (66 $\rightarrow$ 88) and its bodies are unstable across retrains; it is not a usable offline source. (iii) On the long-trained body the heads sit on or just under the NW re-solve under both lags, the RKHS head best; the stale body alone remains useless after the spot move. (iv) The RKHS head of the first pass had no positivity guard and a fixed kernel width, and drove the denominator to its floor on the Feller-violating scenarios (s09: 858 bp on the long-trained body); the guard and the adaptive width of Section 4.4 were introduced here, and the suite’s RKHS rows were recomputed with them.

5.3 The online budget

Setup. The long-trained explicit body at 500 000 particles with each head, against NW re-solved from scratch, at online budgets of 10 000, 30 000 and 80 000 particles (every fit on the whole cloud), both lags, six scenarios, two seeds; the 100 000 column is Table 12. Experiment `suite_budget`.

Conclusions. At 100 000 particles every method sits within a few basis points of the floor and the comparison cannot separate them. As the online cloud thins, NW degrades (49 $\rightarrow$ 57 bp from 100k to 10k) and the body with a head does not (47 $\rightarrow$ 50): at 10 000 particles the RKHS head is 7 bp ahead under both lags, 9 bp in the wings and 6 bp at two years, and the body with either kernel head at 10k matches NW at 80k, with eight times fewer particles and a seventh of the online time. The head is doing the work, not the body: the stale body alone is flat at 61 bp and useless after the spot move at every budget. The ridge head is the weakest at 10k, where it refits 65 features on the thinnest data.

5.4 What the studies say

With the variance floor in place, the classical one-pass kernels are the best cold calibrators and sit a few basis points above the scheme’s own floor, so at 100 000 particles no estimator can look better than “a few basis points”. A per-slice network is only competitive once it is trained properly, and then it is the best cold calibrator. A frozen, properly trained body with an online head reprices today’s vanillas at least as well as a fresh NW calibration at any online budget, and clearly better when the online budget is small, which is where the latency argument for the head lives. The implicit loop is not a useful offline source. The remaining distance to a 10 bp tolerance on liquid quotes is the scheme’s, not the estimator’s.

Table 7: Lagged suite, SSVI scenarios, the same runs in price space: absolute option-price error in basis points of spot (spot is one), mean over scenarios and seeds. Relative price error is not shown; in the far wings the target prices are of order 10^{-6} of spot.

method	offline N	pooled	near-money	wings	worst quote
<i>Surface lag</i>					
Stale leverage, no action	—	16.8	26.9	7.7	61.0
NW re-solve on S_1	—	3.0	4.4	1.9	11.0
Explicit NN, stale f , fresh Dupire	200k	4.9	7.2	2.8	16.0
Explicit NN, stale f , fresh Dupire	500k	4.9	7.2	2.8	16.1
Implicit NN, stale f , fresh Dupire	200k	4.9	7.3	2.8	17.5
Implicit NN, stale f , fresh Dupire	500k	5.8	8.8	3.2	20.6
Explicit NN + RKHS head	200k	2.8	3.9	1.8	9.6
Explicit NN + RKHS head	500k	2.8	3.9	1.8	9.7
Implicit NN + RKHS head	200k	2.8	3.9	1.8	9.7
Implicit NN + RKHS head	500k	2.8	3.9	1.7	9.7
Explicit NN + ridge head	200k	2.7	3.6	1.8	9.4
Explicit NN + ridge head	500k	2.6	3.5	1.8	9.1
Implicit NN + ridge head	200k	2.9	4.0	1.9	10.1
Implicit NN + ridge head	500k	2.8	3.9	1.8	10.2
Explicit NN + spline head	200k	2.7	3.9	1.7	9.7
Explicit NN + spline head	500k	2.7	3.9	1.7	9.7
Implicit NN + spline head	200k	2.8	3.9	1.8	9.7
Implicit NN + spline head	500k	2.7	3.9	1.7	9.7
<i>Surface plus 2% sticky-strike spot lag</i>					
Stale leverage, no action	—	13.7	21.5	6.5	52.4
NW re-solve on S_1	—	3.1	4.5	1.9	11.2
Explicit NN, stale f , fresh Dupire	200k	10.5	16.6	5.2	30.7
Explicit NN, stale f , fresh Dupire	500k	10.5	16.6	5.2	30.5
Implicit NN, stale f , fresh Dupire	200k	10.1	15.8	5.0	29.9
Implicit NN, stale f , fresh Dupire	500k	10.7	16.9	5.4	33.1
Explicit NN + RKHS head	200k	2.8	3.9	1.8	9.7
Explicit NN + RKHS head	500k	2.8	3.9	1.8	9.7
Implicit NN + RKHS head	200k	2.8	3.9	1.9	9.8
Implicit NN + RKHS head	500k	2.9	3.9	1.9	10.7
Explicit NN + ridge head	200k	2.9	4.1	1.8	10.9
Explicit NN + ridge head	500k	2.9	4.2	1.8	11.1
Implicit NN + ridge head	200k	3.1	4.4	1.9	11.6
Implicit NN + ridge head	500k	3.0	4.2	1.9	11.3
Explicit NN + spline head	200k	2.8	3.9	1.8	9.8
Explicit NN + spline head	500k	2.8	3.9	1.8	9.8
Implicit NN + spline head	200k	2.9	3.9	2.0	10.4
Implicit NN + spline head	500k	3.1	4.1	2.2	11.4

Table 8: Lagged suite, Heston-market scenarios..

method	offline N	pooled MAE	pooled RMSE	wings MAE	online latency (s)
<i>Surface lag</i>					
Stale leverage, no action	–	297	323	210	0.0
NW re-solve on S_1	–	233	258	160	8.6
Explicit NN, stale f , fresh Dupire	200k	250	273	174	0.0
Explicit NN, stale f , fresh Dupire	500k	247	270	170	0.0
Implicit NN, stale f , fresh Dupire	200k	233	260	155	0.0
Implicit NN, stale f , fresh Dupire	500k	236	262	158	0.0
Explicit NN + RKHS head	200k	234	258	161	17.1
Explicit NN + RKHS head	500k	234	258	161	17.1
Implicit NN + RKHS head	200k	232	257	158	17.2
Implicit NN + RKHS head	500k	232	257	157	17.1
Explicit NN + ridge head	200k	232	256	160	28.7
Explicit NN + ridge head	500k	232	256	159	27.7
Implicit NN + ridge head	200k	229	254	154	36.4
Implicit NN + ridge head	500k	227	252	155	33.9
Explicit NN + spline head	200k	234	258	161	7.7
Explicit NN + spline head	500k	234	258	161	7.7
Implicit NN + spline head	200k	233	258	160	7.3
Implicit NN + spline head	500k	232	257	159	7.4
<i>Surface plus 2% sticky-strike spot lag</i>					
Stale leverage, no action	–	281	313	195	0.0
NW re-solve on S_1	–	225	248	154	8.5
Explicit NN, stale f , fresh Dupire	200k	227	253	155	0.0
Explicit NN, stale f , fresh Dupire	500k	225	251	152	0.0
Implicit NN, stale f , fresh Dupire	200k	217	245	140	0.0
Implicit NN, stale f , fresh Dupire	500k	214	241	138	0.0
Explicit NN + RKHS head	200k	225	248	154	17.2
Explicit NN + RKHS head	500k	225	248	154	17.3
Implicit NN + RKHS head	200k	224	247	151	17.3
Implicit NN + RKHS head	500k	224	247	150	17.2
Explicit NN + ridge head	200k	221	244	150	30.7
Explicit NN + ridge head	500k	221	244	150	30.4
Implicit NN + ridge head	200k	218	243	145	38.8
Implicit NN + ridge head	500k	218	242	146	39.0
Explicit NN + spline head	200k	226	249	155	7.6
Explicit NN + spline head	500k	226	249	155	7.6
Implicit NN + spline head	200k	226	250	154	7.3
Implicit NN + spline head	500k	225	248	153	7.3

Table 9: Lagged suite, Heston-market scenarios, the same runs in price space: absolute option-price error in basis points of spot (spot is one), mean over scenarios and seeds. Relative price error is not shown; in the far wings the target prices are of order 10^{-6} of spot.

method	offline N	pooled	near-money	wings	worst quote
<i>Surface lag</i>					
Stale leverage, no action	—	75.9	132.4	27.5	186.9
NW re-solve on S_1	—	60.3	105.3	21.8	155.0
Explicit NN, stale f , fresh Dupire	200k	64.3	111.8	23.5	160.8
Explicit NN, stale f , fresh Dupire	500k	63.7	111.2	23.1	160.1
Implicit NN, stale f , fresh Dupire	200k	61.5	107.7	22.0	157.6
Implicit NN, stale f , fresh Dupire	500k	62.1	108.3	22.5	157.5
Explicit NN + RKHS head	200k	60.6	105.5	22.0	155.2
Explicit NN + RKHS head	500k	60.5	105.5	22.0	155.1
Implicit NN + RKHS head	200k	60.5	105.4	21.9	155.0
Implicit NN + RKHS head	500k	60.5	105.5	21.9	155.2
Explicit NN + ridge head	200k	60.1	104.6	22.0	154.6
Explicit NN + ridge head	500k	60.0	104.5	21.9	154.3
Implicit NN + ridge head	200k	60.2	104.9	21.9	155.1
Implicit NN + ridge head	500k	59.6	103.5	22.0	153.7
Explicit NN + spline head	200k	60.5	105.4	22.0	155.2
Explicit NN + spline head	500k	60.5	105.4	22.0	155.2
Implicit NN + spline head	200k	60.5	105.4	22.0	155.2
Implicit NN + spline head	500k	60.5	105.4	22.0	155.2
<i>Surface plus 2% sticky-strike spot lag</i>					
Stale leverage, no action	—	70.9	126.0	23.6	179.8
NW re-solve on S_1	—	58.5	102.0	21.1	152.2
Explicit NN, stale f , fresh Dupire	200k	58.1	103.2	19.4	148.7
Explicit NN, stale f , fresh Dupire	500k	57.7	102.6	19.1	147.8
Implicit NN, stale f , fresh Dupire	200k	57.3	102.3	18.8	147.8
Implicit NN, stale f , fresh Dupire	500k	56.8	101.1	18.8	145.7
Explicit NN + RKHS head	200k	58.6	102.3	21.2	152.4
Explicit NN + RKHS head	500k	58.6	102.2	21.2	152.3
Implicit NN + RKHS head	200k	58.7	102.4	21.2	152.5
Implicit NN + RKHS head	500k	58.7	102.4	21.2	152.7
Explicit NN + ridge head	200k	57.6	100.5	20.9	150.4
Explicit NN + ridge head	500k	57.6	100.3	20.9	150.3
Implicit NN + ridge head	200k	58.1	101.3	21.0	152.0
Implicit NN + ridge head	500k	57.7	100.4	21.1	150.9
Explicit NN + spline head	200k	58.8	102.4	21.4	152.8
Explicit NN + spline head	500k	58.8	102.5	21.5	152.8
Implicit NN + spline head	200k	58.9	102.4	21.5	152.8
Implicit NN + spline head	500k	58.8	102.4	21.4	152.8

Table 10: Relative error of the fitted denominator against the exact conditional expectation, in percent, and seconds per fit.

estimator	s02 wings		s11 wings		s11 centre	
	$t = 0.25$	$t = 1$	$t = 0.25$	$t = 1$	mean	s / fit
NW, 30k particles	3.9	6.0	17.5	10.3	14.0	0
NW, 100k	2.8	3.5	14.6	9.5	15.0	0
NW, 500k	1.9	1.6	14.9	6.1	13.7	1
Spline, 100k	4.6	4.1	15.4	12.6	13.1	0
RKHS, 100k	4.1	9.6	29.8	12.7	34.7	0
Net 64×2 , 120 full-batch steps (paper)	6.2	2.6	55.1	16.2	33.8	6
Net 64×2 , 1000 full-batch steps	2.2	1.3	54.6	10.0	28.7	52
Net 64×2 , 20000 minibatch steps	0.6	1.7	12.7	9.0	14.9	145
Net 64×3 , 120 full-batch steps	1.9	1.6	55.3	22.5	30.0	9
Net 64×3 , 20000 minibatch steps	0.7	1.8	10.4	9.7	15.2	230
Net 256×2 , 20000 minibatch steps	1.0	1.5	8.0	10.5	15.3	520

Table 11: Offline bodies at 500 000 particles as cold calibrators, six scenarios, seed 0: pooled and wings MAE (bp), the two-year slice, and training time.

training	body	pooled MAE	wings MAE	$T = 2$ MAE	train s
Short training, 30k fits	explicit	64	93	17	346
Short training, 30k fits	implicit	66	96	28	756
Short training, 100k fits	explicit	70	102	18	1262
Short training, 100k fits	implicit	78	115	21	1712
Long training, 100k fits	explicit	50	72	18	9115
Long training, 100k fits	implicit	88	131	26	11671

Table 12: Lagged suite on the 500 000-particle explicit body, six scenarios, two online seeds, pooled MAE (bp), by body training and fit sample. Heads are the guarded versions of Section 4.4.

online method	surface lag			surface plus spot lag		
	short, 30k fits	short, 100k fits	long, 100k fits	short, 30k fits	short, 100k fits	long, 100k fits
NW re-solve	52	49	49	52	48	48
Explicit body, stale f , fresh Dupire	64	65	61	119	118	99
Explicit body + RKHS head	53	56	47	53	55	45
Explicit body + spline head	52	48	48	51	48	49
Explicit body + ridge head	63	59	48	57	60	47

Table 13: Pooled MAE (bp) by online particle budget, and the online time at 10 000 particles.

online method	surface lag, pooled MAE				surface plus spot lag				s at 10k
	10k	30k	80k	100k	10k	30k	80k	100k	
NW re-solve	57	52	50	49	56	51	49	48	1.0
Explicit body, stale f , fresh Dupire	61	61	61	61	99	99	99	99	0.0
Explicit body + spline head	51	50	48	48	51	49	49	49	0.9
Explicit body + RKHS head	50	48	47	47	49	48	47	45	2.1
Explicit body + ridge head	54	48	49	48	56	49	48	47	2.9